



County Keras network: observed versus predicted

Predictions from a held-out 15% assessment set




County Keras learning curves

Early stopping restores the epoch with minimum validation loss



County Keras permutation importance


Increase in assessment MSE after shuffling each predictor



Importance measures predictive reliance, not causality.

County Keras model performance

Lower RMSE and MAE are better; higher R-squared is better



County Keras residual distributions

Residual = observed minus predicted recovery ratio

